

EMPLOYEE ATTRITION CLASSIFICATION DATASET

An In-Depth Synthetic Simulation for Attrition
Analysis and Prediction

ABSTRACT

This report details the process of deploying a machine learning model using AWS SageMaker, using the Employee Attrition Classification dataset from Kaggle. The dataset which includes various features related to employee demographics, job roles, serves as a basis for predicting employee turnover within an organization. The report outlines the entire workflow, from data preparation to model training, evaluation, and deployment. Key aspects of AWS are explored, including configuring the SageMaker environment, training the model, and deploying it for inference. Additionally, the project required the management of data and model files, which were stored in an Amazon S3 bucket. The objective is to demonstrate a practical application of AWS SageMaker in implementing machine learning solutions for business problems, showcasing the effectiveness and efficiency of cloud-based machine learning deployments in addressing employee attrition and enhancing organizational retention strategies.

Thapelo Lenzi

Deploying Machine Learning Models in AWS
SageMaker

Dataset

This dataset consists of 74 498 entries with each including a unique Employee ID and features that influence employee attrition. The goal is to understand the factors contributing to attrition (Whether the employee has left the company, encoded as 0 for stayed and 1 for left) and develop predictive models to identify at-risk employees. The features include:

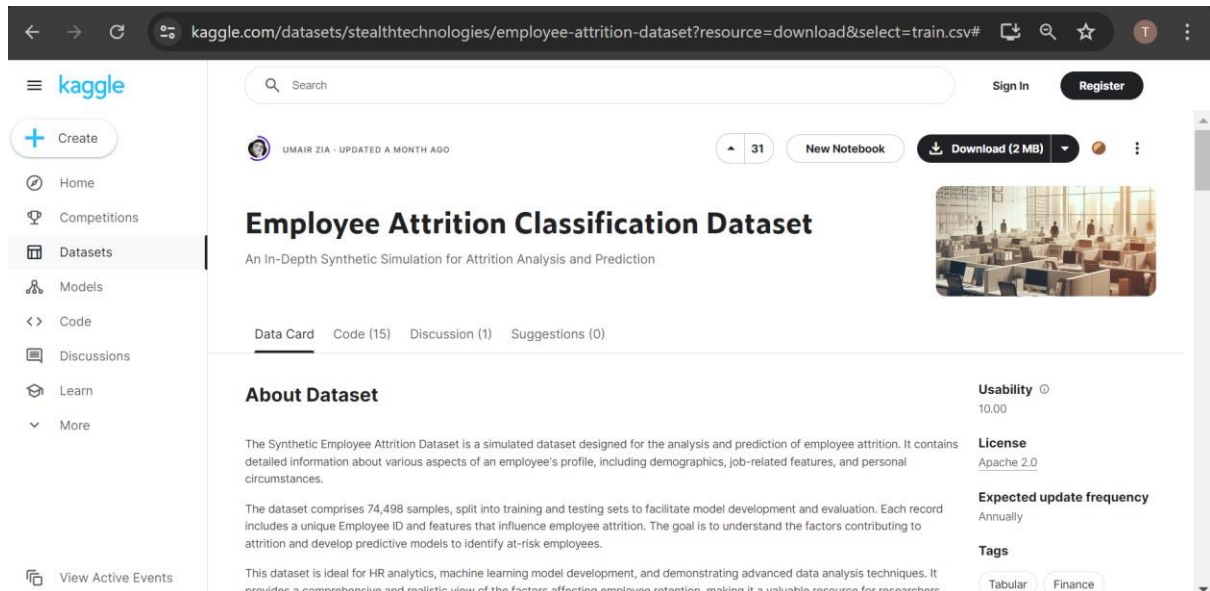
- Age
- Gender
- Years at company
- Monthly income
- Job role
- Work-life balance
- Education level
- And more

<https://www.kaggle.com/datasets/stealthtechnologies/employee-attritiondataset?resource=download&select=train.csv>

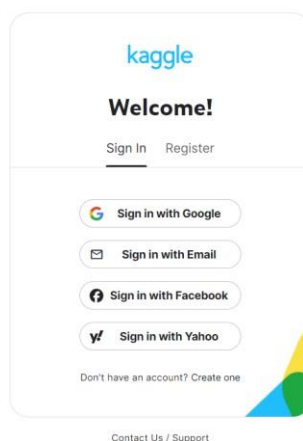
Data Extraction

Obtaining the dataset from Kaggle:

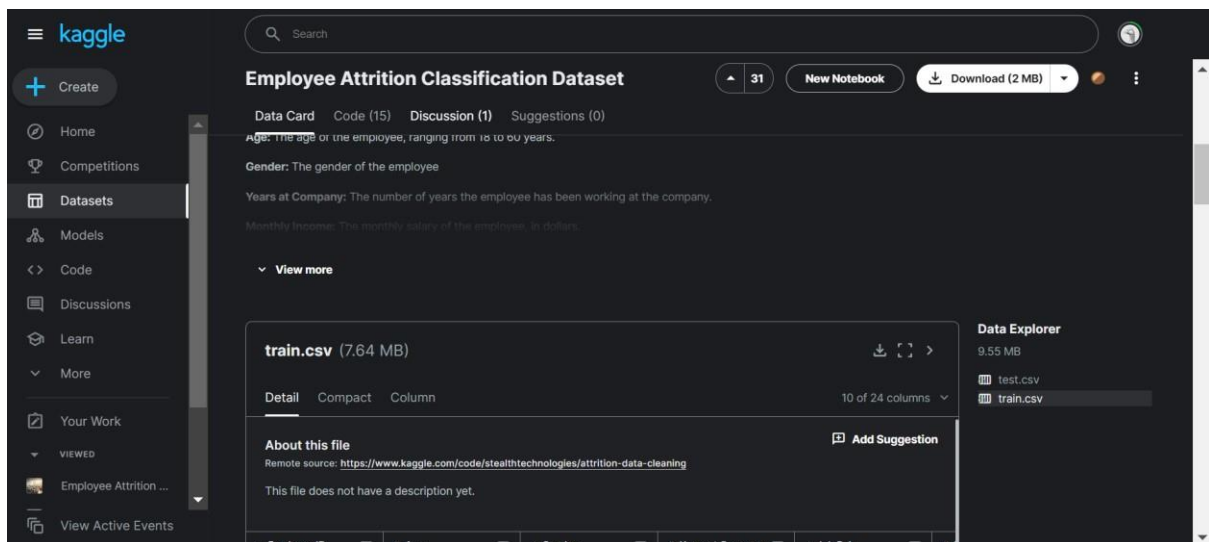
Search for the Employee Attrition Classification dataset on Kaggle.



Sign In:



Go to the dataset page and click on the download button to manually download the dataset:



Data Preparation

The data preparation and exploratory data analysis stages of this project are done using Google Colab.

Loading Libraries, Importing necessary libraries for data manipulation and visualization:

```
Necessary Packages

[20] import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
pd.set_option('display.max_columns', None)
import warnings
warnings.filterwarnings('ignore')

from sklearn.preprocessing import LabelEncoder
```

Loading the data into a pandas data frame:

```
[4] #Load the data
data = pd.read_csv("/content/Attrition.csv")
```

Exploratory Data Analysis

The overview of the dataset:

```
[5] #Display the top 5 rows of the dataset
data.head()
```

	Employee ID	Age	Gender	Years at Company	Job Role	Monthly Income	Work-Life Balance	Job Satisfaction	Performance Rating	Number of Promotions	Overtime	Distance from Home	Education Level
0	8410	31	Male	19	Education	5390	Excellent	Medium	Average	2	No	22	Associate Degree
1	64756	59	Female	4	Media	5534	Poor	High	Low	3	No	21	Master's Degree
2	30257	24	Female	10	Healthcare	8159	Good	High	Low	0	No	11	Bachelor's Degree
3	65791	36	Female	7	Education	3989	Good	High	High	1	No	27	High School
4	65026	56	Male	41	Education	4821	Fair	Very High	Average	0	Yes	71	High School

```
#Display the number of rows and columns
data.shape
```

```
(8185, 24)
```

```
#Data type summary
data.info()
```

RangeIndex: 8186 entries, 0 to 8185			
Data columns (total 24 columns):			
#	Column	Non-Null Count	Dtype
0	Employee ID	8186 non-null	int64
1	Age	8186 non-null	int64
2	Gender	8186 non-null	object
3	Years at Company	8186 non-null	int64
4	Job Role	8186 non-null	object
5	Monthly Income	8186 non-null	int64
6	Work-Life Balance	8186 non-null	object
7	Job Satisfaction	8186 non-null	object
8	Performance Rating	8186 non-null	object
9	Number of Promotions	8186 non-null	int64
10	Overtime	8186 non-null	object
11	Distance from Home	8186 non-null	int64
12	Education Level	8186 non-null	object
13	Marital Status	8186 non-null	object
14	Number of Dependents	8186 non-null	int64
15	Job Level	8186 non-null	object
16	Company Size	8186 non-null	object
17	Company Tenure	8186 non-null	int64
18	Remote Work	8186 non-null	object
19	Leadership Opportunities	8186 non-null	object
20	Innovation Opportunities	8186 non-null	object
21	Company Reputation	8185 non-null	object
22	Employee Recognition	8185 non-null	object
23	Attrition	8185 non-null	object
dtypes: int64(8), object(16)			

```
[7] #Statistical summary
data.describe()
```

	Employee ID	Age	Years at Company	Monthly Income	Number of Promotions	Distance from Home	Number of Dependents	Company Tenure
count	8186.000000	8186.000000	8186.000000	8186.000000	8186.000000	8186.000000	8186.000000	8186.000000
mean	37287.220743	38.615075	15.667237	7336.613609	0.827144	49.724408	1.649890	55.863059
std	21446.665299	12.135169	11.291536	2141.167892	0.996414	28.518282	1.556062	25.546607
min	8.000000	18.000000	1.000000	1855.000000	0.000000	1.000000	0.000000	2.000000
25%	18757.250000	28.000000	7.000000	5714.000000	0.000000	25.000000	0.000000	36.000000
50%	37088.500000	39.000000	13.000000	7373.000000	0.000000	50.000000	1.000000	56.000000
75%	55937.750000	49.000000	23.000000	8879.500000	1.000000	74.000000	3.000000	76.000000
max	74488.000000	59.000000	51.000000	15495.000000	4.000000	99.000000	6.000000	127.000000

The above snippets provide a sample of the data and understand the structure of the dataset, including column names and initial values. They provide the dimensions of the dataset, which can be observed as 8185 rows and 24 columns, which help in understanding the size of the dataset.

The data type and statistical summaries, include the number of non-null entries in each column and datatypes together with metrics such as the mean, minimum and maximum. These help with understanding the distribution and range of the numeric data and identifying any missing values. From the above, it can be noted that the dataset has two data types (int64 and object) and the last three columns have some null values.

Handling Missing Data

The missing values can be explicitly displayed, as shown below. Since there is only one value missing in each of the three columns and there are a lot of entries, the missing values can be dropped and will not make any significant change in the analysis.

```
[11] #Get number of duplicated data  
data.duplicated().sum()
```

```
0
```

```
[8] #Get number of empty rows in each column  
data.isnull().sum()
```

```
Employee ID      0  
Age              0  
Gender           0  
Years at Company 0  
Job Role         0  
Monthly Income   0  
Work-Life Balance 0  
Job Satisfaction 0  
Performance Rating 0  
Number of Promotions 0  
Overtime         0  
Distance from Home 0  
Education Level  0  
Marital Status   0  
Number of Dependents 0  
Job Level        0  
Company Size     0  
Company Tenure   0  
Remote Work      0  
Leadership Opportunities 0  
Innovation Opportunities 0  
Company Reputation 1  
Employee Recognition 1  
Attrition        1  
dtype: int64
```

```
[12] #Drop the empty values  
data.dropna(inplace = True)
```

```
[14] #Verify the empty values are removed  
data.isnull().values.any()
```

```
False
```

```
[19] #Display the count summary for each column
for column in data.columns:
    if data[column].dtype == object:
        print(f"{column}: {data[column].unique()}")
        print(data[column].value_counts())
        print('.....')

[19] .....
Innovation Opportunities: ['No' 'Yes']
Innovation Opportunities
No      6873
Yes     1312
Name: count, dtype: int64
.....
Company Reputation: ['Excellent' 'Fair' 'Poor' 'Good']
Company Reputation
Good      4036
Poor     1667
Fair     1655
Excellent   827
Name: count, dtype: int64
.....
Employee Recognition: ['Medium' 'Low' 'High' 'Very High']
Employee Recognition
Low       3267
Medium   2439
High     2062
Very High   417
Name: count, dtype: int64
.....
Attrition: ['Stayed' 'Left']
Attrition
Stayed   4282
Left     3983
Name: count, dtype: int64
.....
```

The above is a summary of each column. In the dependent column (Attrition), there are two unique values with 'Stayed' having the greater count.

Feature Engineering

The dependent column is originally in text form. In order for it to be used in the model, it is converted to numerical form using the LabelEncoder from the scikit-learn library, as seen in the snippet below.

```
le = LabelEncoder()
data['Attrition'] = le.fit_transform(data['Attrition'])
data.head(5)
```

	Number of Promotions	Overtime	Distance from Home	Education Level	Marital Status	Number of Dependents	Job Level	Company Size	Company Tenure	Remote Work	Leadership Opportunities	Innovation Opportunities	Company Reputation	Employee Recognition	Attrition
3	2	No	22	Associate Degree	Married	0	Mid	Medium	89	No	No	No	Excellent	Medium	1
4	3	No	21	Master's Degree	Divorced	3	Mid	Medium	21	No	No	No	Fair	Low	1
4	0	No	11	Bachelor's Degree	Married	3	Mid	Medium	74	No	No	No	Poor	Low	1
1	1	No	27	High School	Single	2	Mid	Small	50	Yes	No	No	Good	Medium	1
3	0	Yes	71	High School	Divorced	0	Senior	Medium	68	No	No	No	Fair	Medium	1

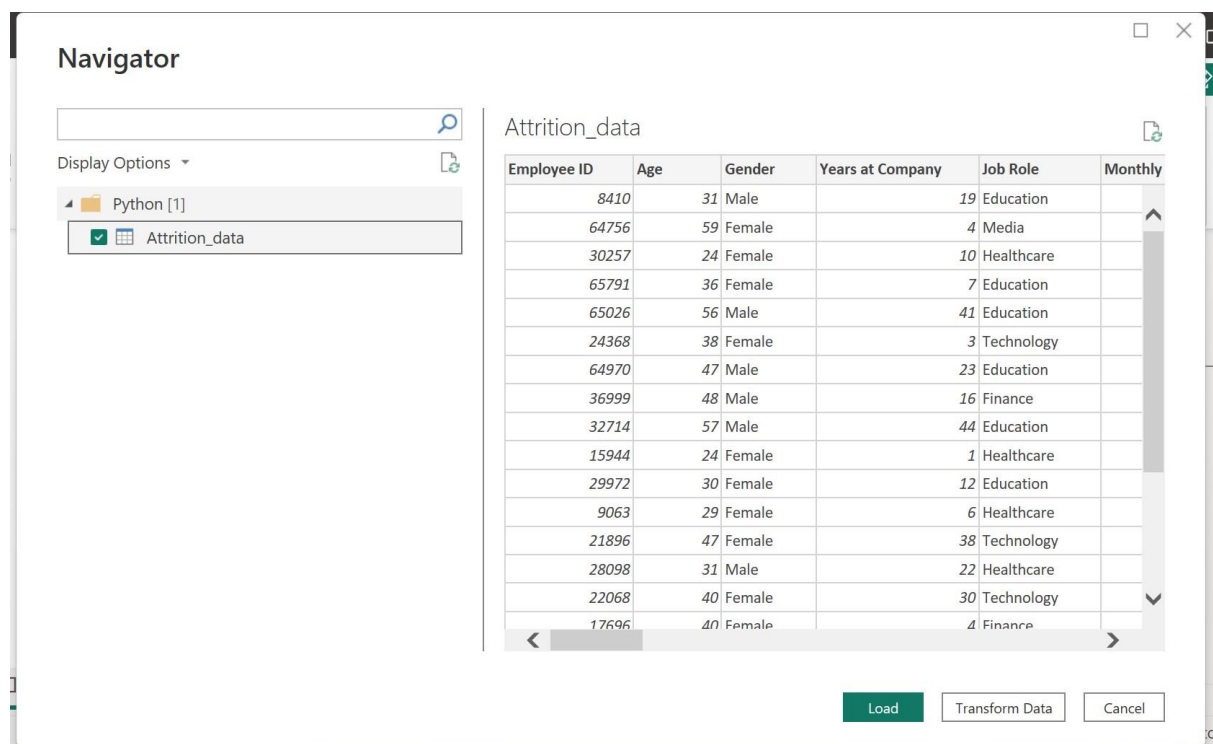
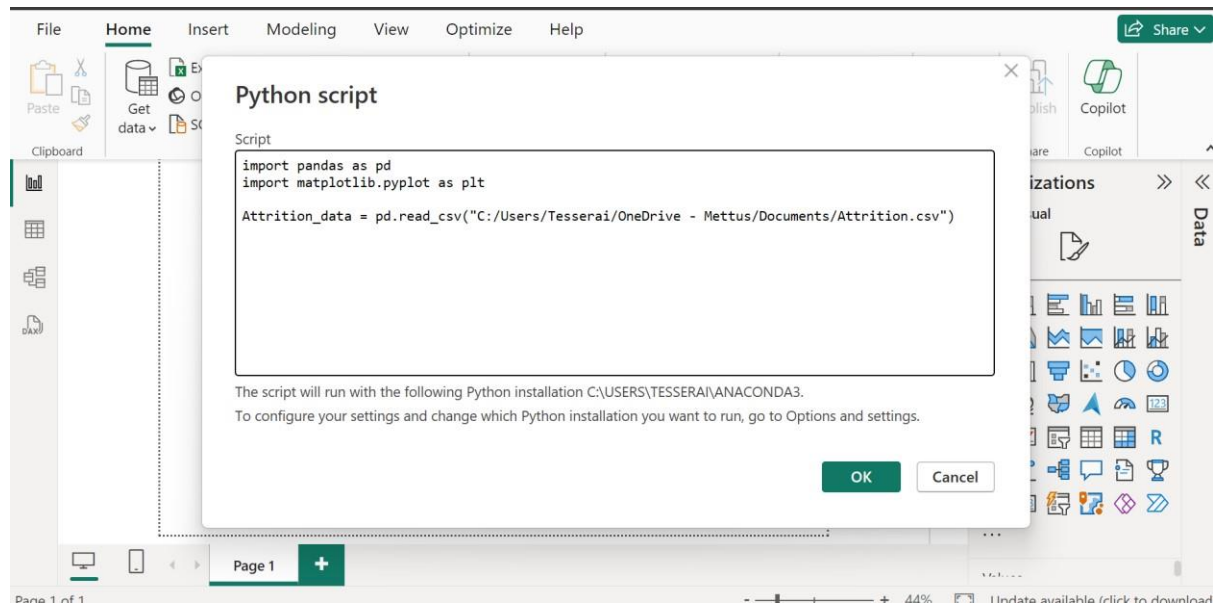
```
le.classes_
array(['Left', 'Stayed'], dtype=object)

[23] Employee_Attrition = le.classes_
print(Employee_Attrition)

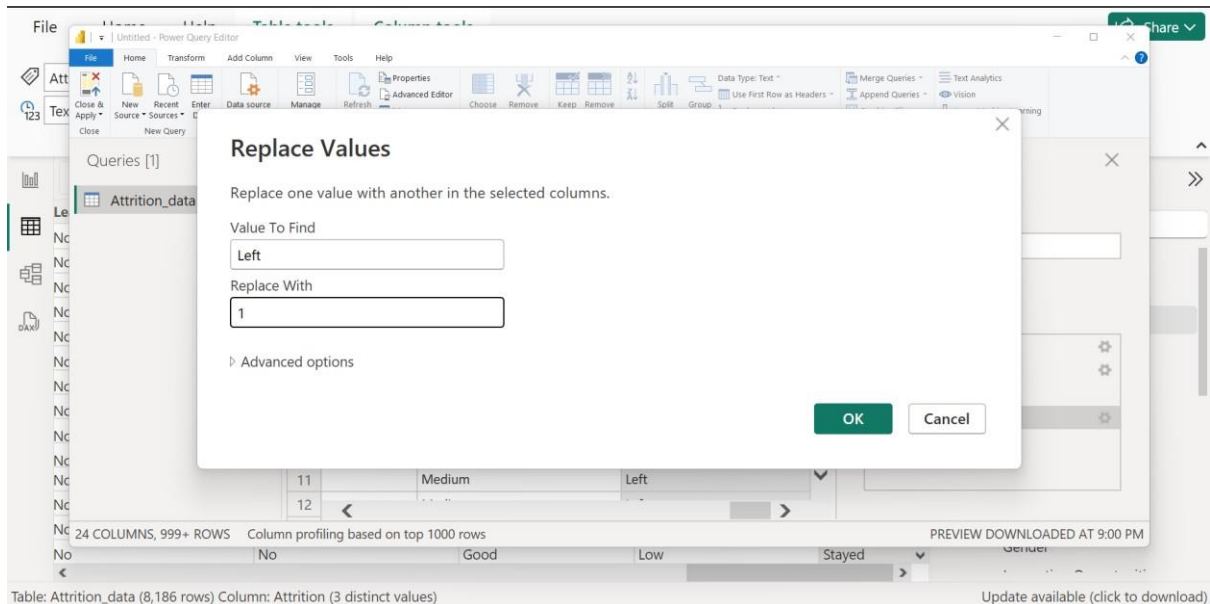
['Left' 'Stayed']
```

Data Visualization using PowerBI Desktop

Connecting to Python Script in PowerBI Desktop:

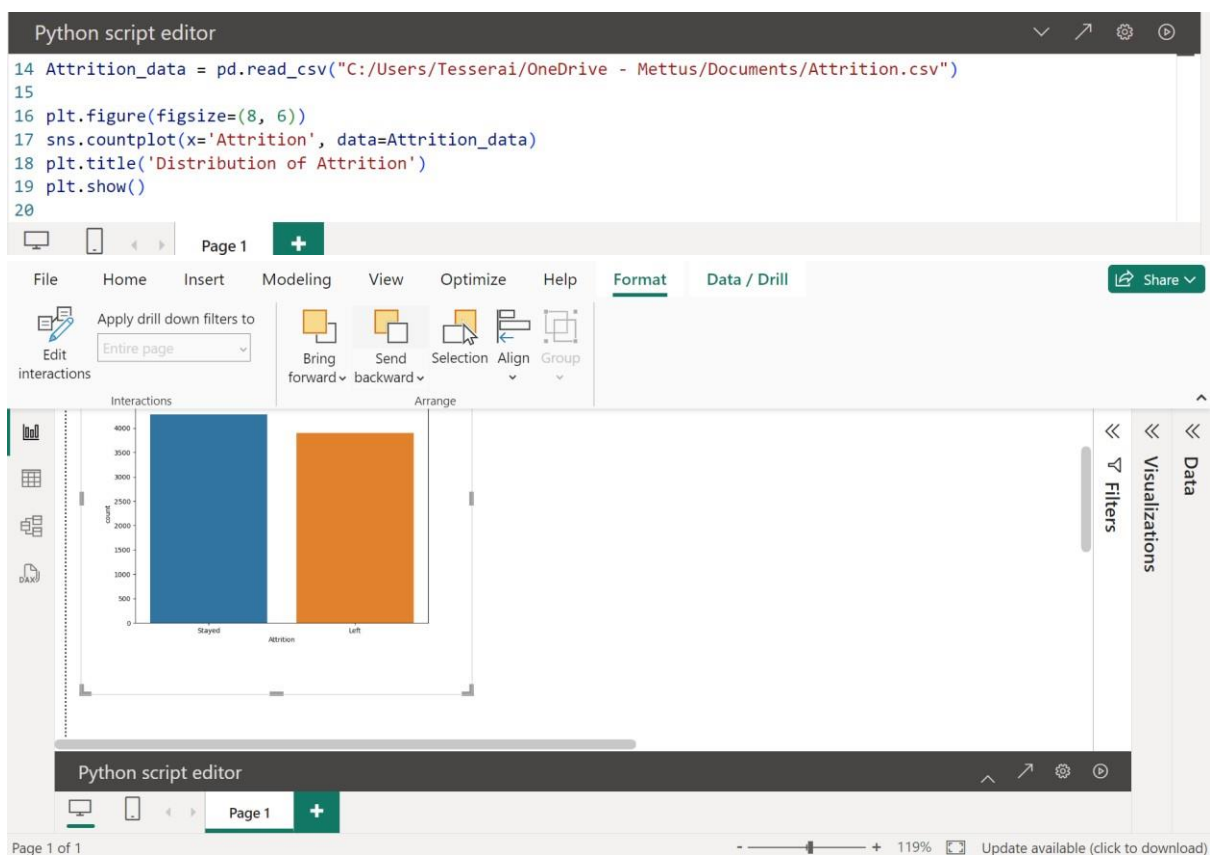


Transforming the Attrition data column from categorical to numerical:



The screenshot shows the Power Query Editor interface. A 'Replace Values' dialog box is open, allowing the user to replace one value with another in the selected columns. The 'Value To Find' field contains 'Left', and the 'Replace With' field contains '1'. The dialog also has an 'Advanced options' section and 'OK' and 'Cancel' buttons. In the background, a table is visible with columns for 'Attrition', 'Medium', and 'Left'. The table has 24 columns and 999+ rows. The status bar at the bottom indicates 'Table: Attrition_data (8,186 rows) Column: Attrition (3 distinct values)'.

Attrition	Medium	Left
No	Good	Low
Stayed	Good	Low

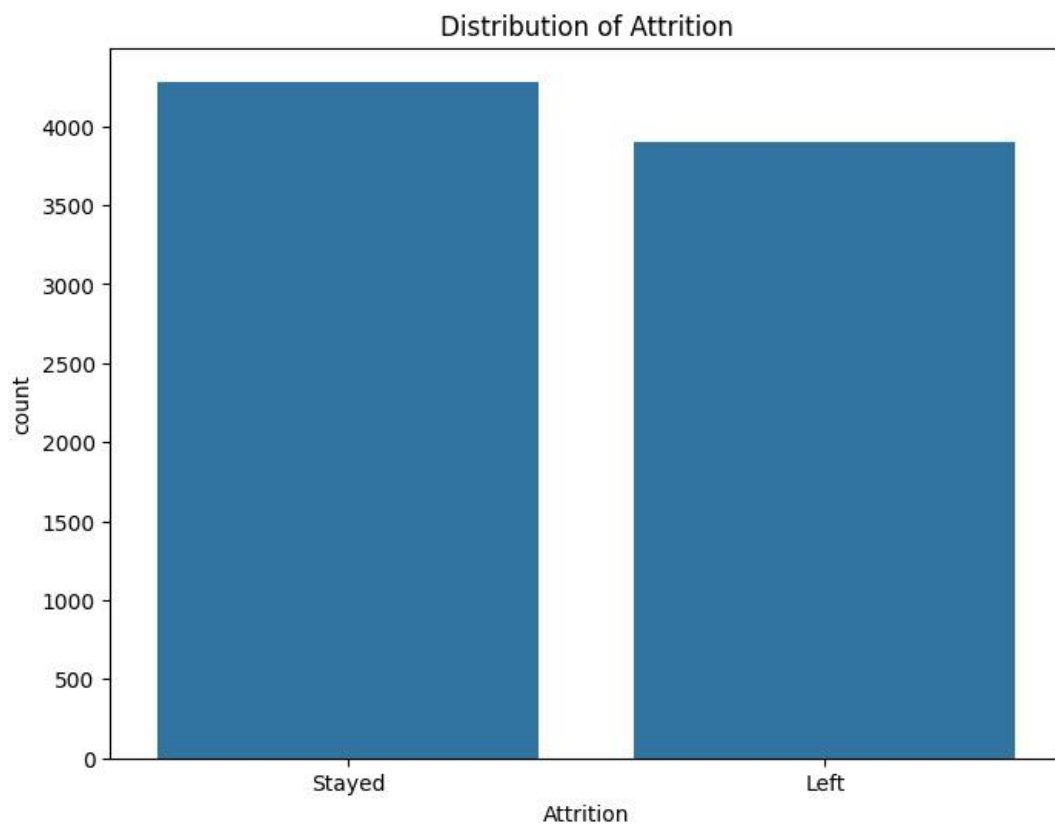



The screenshot shows a Python script editor and a bar chart. The script reads a CSV file and creates a countplot of the 'Attrition' column. The bar chart shows the distribution of 'Attrition' with 'Stayed' and 'Left' categories. The status bar at the bottom indicates 'Page 1 of 1' and '119%' zoom.


```
14 Attrition_data = pd.read_csv("C:/Users/Tessera/OneDrive - Mettus/Documents/Attrition.csv")
15
16 plt.figure(figsize=(8, 6))
17 sns.countplot(x='Attrition', data=Attrition_data)
18 plt.title('Distribution of Attrition')
19 plt.show()
20
```

Attrition	count
Stayed	3500
Left	3000

More EDA Visuals







Data Splitting

```
[29] #Libraries
      from sklearn.model_selection import train_test_split

[30] x = data.drop('Attrition', axis=1)
      y = data['Attrition']

      x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.2, random_state=42)

[31] #Shape of the train and test data
      x_train.shape
      (6548, 23)

[32] y_train.shape
      (6548,)

[33] x_test.shape
      (1637, 23)

[34] y_test.shape
      (1637,)
```

Model Training

The primary goal of the dataset is to classify employees into two categories: those who stayed and those who left the company. Logistic regression is specifically designed for binary classification tasks, making it ideal for distinguishing between these two outcomes. It estimates the probability that an employee falls into one of the two categories

```
[25] #Libraries
      from sklearn.linear_model import LogisticRegression
      from sklearn import metrics

[26] model = LogisticRegression()

[27] model.fit(x_train, y_train)

LogisticRegression
LogisticRegression()

[28] expected = y_train
      predicted = model.predict(x_train)
```

Model Evaluation

```
print(metrics.classification_report(expected, predicted))
```

	precision	recall	f1-score	support
0	0.66	0.65	0.65	3126
1	0.68	0.70	0.69	3422
accuracy			0.67	6548
macro avg	0.67	0.67	0.67	6548
weighted avg	0.67	0.67	0.67	6548

```
[30] print(metrics.confusion_matrix(expected, predicted))
```

[[2020 1106]
[1039 2383]]

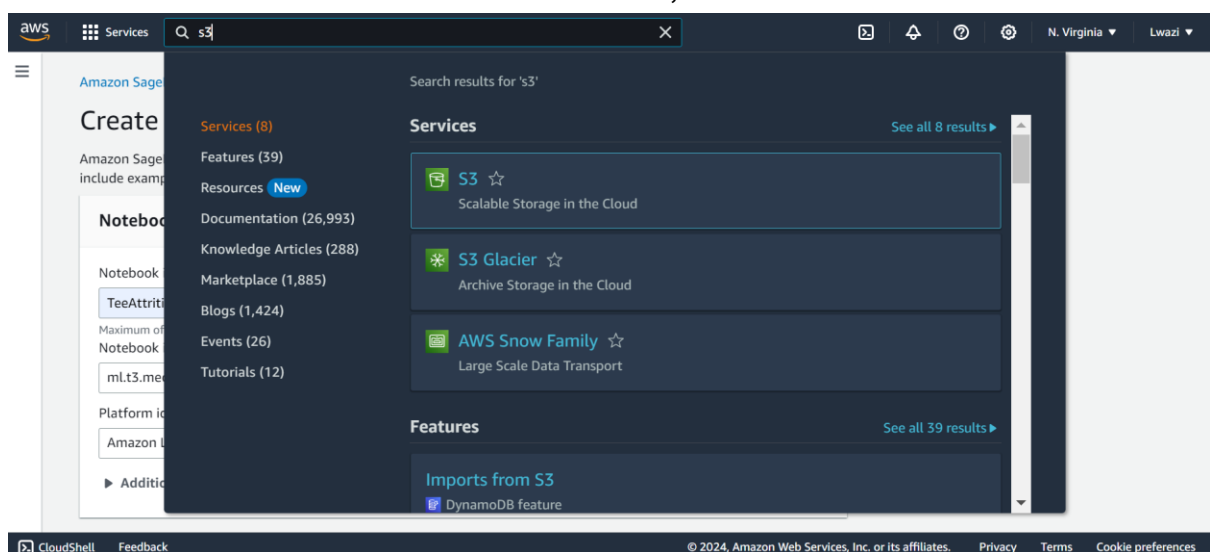
From the above metrics snippets, the model shows a balanced performance in classifying between the two classes. For the '0', the precision is 66% and the recall is 65%, indicating a moderate performance in identifying the '0' class. The '1' class, however, there is a better performance in detecting this class, with a precision and recall of 68% and 70%, respectively. From this, it can be observed that the model has more difficulty predicting the stayed class compared to the left class.

Amazon Web Service

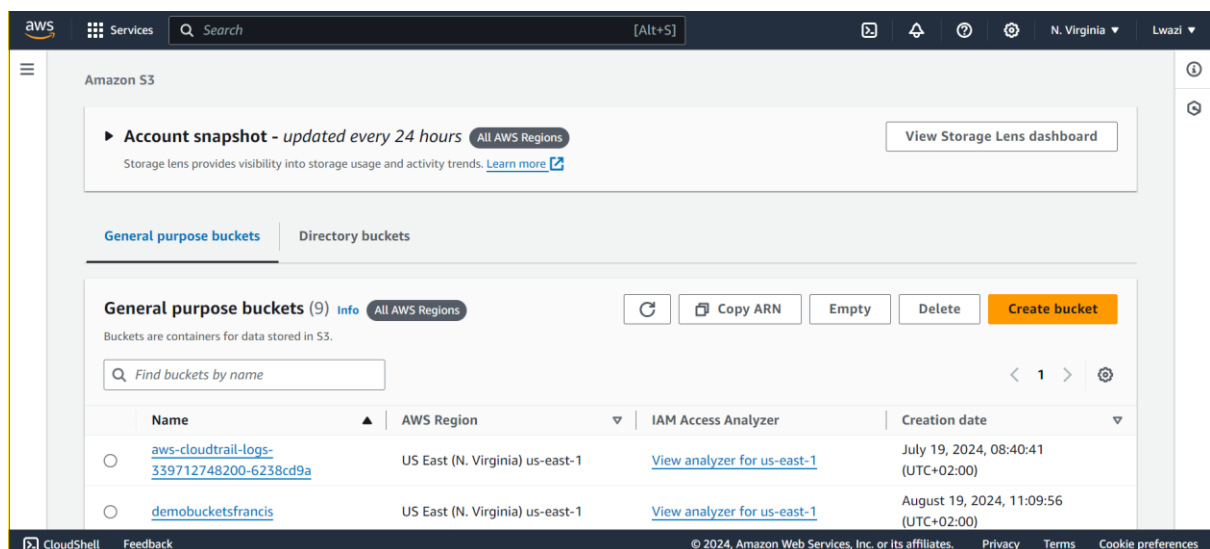
Creating an S3 Bucket:

An Amazon S3 bucket named 'sagemakerteebucket' is created to store the datasets and model files used in the project. The configuration is set to general purpose, which is recommended for most use cases and access patterns. This type of bucket allows for the distribution of storage across multiple availability zones, which ensures redundancy and reliability. The bucket is created in the US East N. Virginia) region. The bucket will be used for managing the input and output data for the AWS SageMaker model training and deployment process.

On the search bar in the aws console dashboard, search for S3:



Create a new bucket:



aws

Services

Search

[Alt+S]

N. Virginia

Lwazi

General configuration

AWS Region

US East (N. Virginia) us-east-1

Bucket type

Info

☒ General purpose

Recommended for most use cases and access patterns. General purpose buckets are the original S3 bucket type. They allow a mix of storage classes that redundantly store objects across multiple Availability Zones.

☐ Directory - New

Recommended for low-latency use cases. These buckets use only the S3 Express One Zone storage class, which provides faster processing of data within a single Availability Zone.

Bucket name

Info

sagemakerteebucket

Bucket name must be unique within the global namespace and follow the bucket naming rules. [See rules for bucket naming](#)

Copy settings from existing bucket - optional

Only the bucket settings in the following configuration are copied.

Choose bucket

Format: s3://bucket/prefix

CloudShell

Feedback

© 2024, Amazon Web Services, Inc. or its affiliates.

Privacy

Terms

Cookie preferences

aws

Services

Search

[Alt+S]

N. Virginia

Lwazi

Successfully created bucket "sagemakerteebucket"

To upload files and folders, or to configure additional bucket settings, choose [View details](#).

General purpose buckets (10)

Info

All AWS Regions

Refresh

Copy ARN

Empty

Delete

Create bucket

Find buckets by name

< 1 >

Name	AWS Region	IAM Access Analyzer	Creation date (UTC+02:00)
<input type="radio"/> nontobekobucket	US East (N. Virginia) us-east-1	View analyzer for us-east-1	August 19, 2024, 09:04:28 (UTC+02:00)
<input type="radio"/> romeodiabetecbucket	US East (N. Virginia) us-east-1	View analyzer for us-east-1	August 19, 2024, 09:03:37 (UTC+02:00)
<input type="radio"/> sagemakeremployeeattrition	US East (N. Virginia) us-east-1	View analyzer for us-east-1	August 18, 2024, 02:26:30 (UTC+02:00)
<input type="radio"/> sagemakerteebucket	US East (N. Virginia) us-east-1	View analyzer for us-east-1	August 19, 2024, 23:21:21 (UTC+02:00)

CloudShell

Feedback

© 2024, Amazon Web Services, Inc. or its affiliates.

Privacy

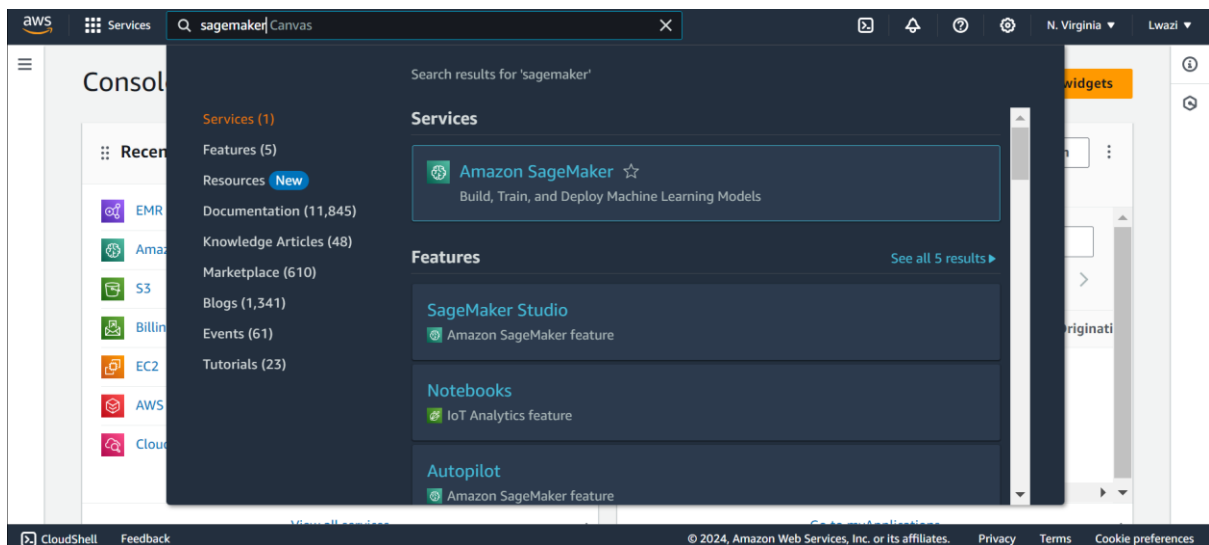
Terms

Cookie preferences

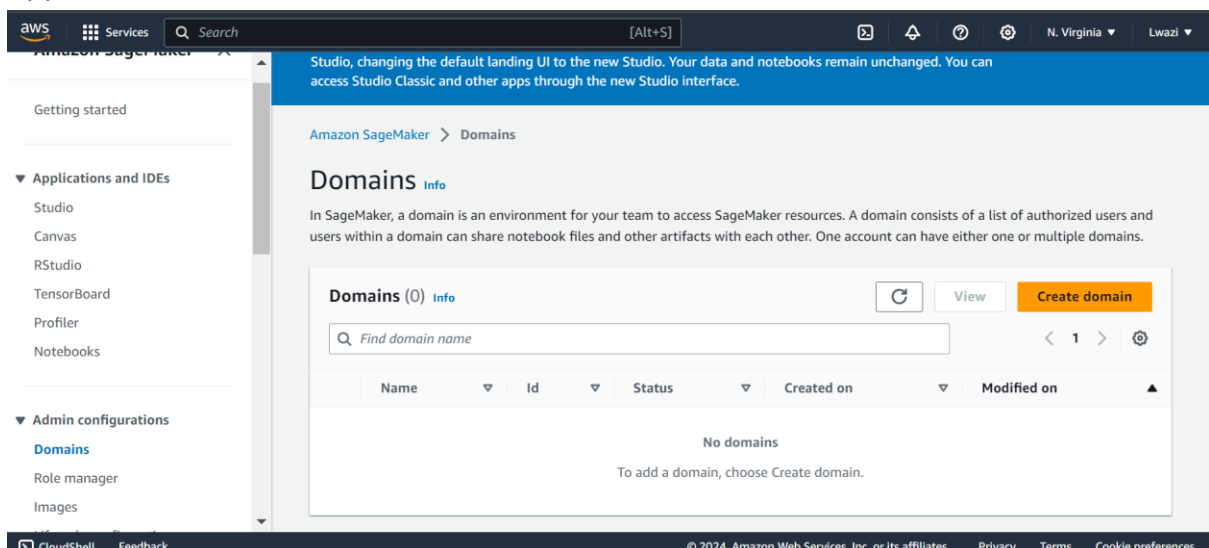
Amazon SageMaker

The following captures the process of setting up a new notebook instance in AWS SageMaker. The instance is named TeeAttrition and is configured to use the 'ml.t3.medium' instance type, which provides a balance of compute, memory and networking resources. This notebook instance is the environment where the model training and evaluation will be conducted.

On the search bar in the aws console dashboard, search for sagemaker:



On the side bar, in the Amazon SageMaker dashboard, navigate to notebooks under Applications and IDEs:



Create a notebook instance:

The screenshot displays the Amazon SageMaker console interface. The top navigation bar includes the AWS logo, 'Services', a search bar, and user information (N. Virginia, Lwazi). The left sidebar shows a menu with 'Getting started', 'Applications and IDEs' (Studio, Canvas, RStudio, TensorBoard, Profiler, Notebooks), and 'Admin configurations' (Domains, Role manager, Images). The main content area is titled 'Notebook instances' and features a 'Create notebook instance' button. Below this is a table of existing instances:

Name	Instance	Creation time	Status	Actions
TsehsDiabetesInstance	ml.t2.medium	8/19/2024, 8:50:59 AM	InService	Open Jupyter Open JupyterLab
Diabetes	ml.t3.medium	8/19/2024, 8:36:54 AM	InService	Open Jupyter Open JupyterLab
		8/19/2024, 8:34:42		Open Jupyter Open

Below the table, the 'Create notebook instance' form is shown with the following fields:

- Notebook instance name:** TeeAttrition
- Notebook instance type:** ml.t3.medium
- Platform identifier:** Amazon Linux 2, Jupyter Lab 3
- Additional configuration:** (expandable section)

The bottom of the console shows the footer with 'CloudShell', 'Feedback', and copyright information for Amazon Web Services, Inc. or its affiliates.

Creating an IAM role:

The IAM role allows the SageMaker notebook instance to access the selected S3 bucket. The role is restricted to a specific bucket, which ensures that the SageMaker instance can read and write only to the designated storage location. This is important for handling input datasets and storing the model outputs.

aws

Services

Search

[Alt+S]

N. Virginia

Lwazi

Permissions and encryption

IAM role

Notebook instances require permissions to call other services including SageMaker and S3. Choose a role or let us create a role with the [AmazonSageMakerFullAccess](#) IAM policy attached.

AmazonSageMaker-ExecutionRole-20240819T085077

▲

Create a new role

▲

Enter a custom IAM role ARN

Use existing role

AmazonSageMaker-ExecutionRole-20240818T022737

AmazonSageMaker-ExecutionRole-20240819T083104

AmazonSageMaker-ExecutionRole-20240819T083116

AmazonSageMaker-ExecutionRole-20240819T083153

AmazonSageMaker-ExecutionRole-20240819T083170

AmazonSageMaker-ExecutionRole-20240819T083176

AmazonSageMaker-ExecutionRole-20240819T083332

▼

CloudShell

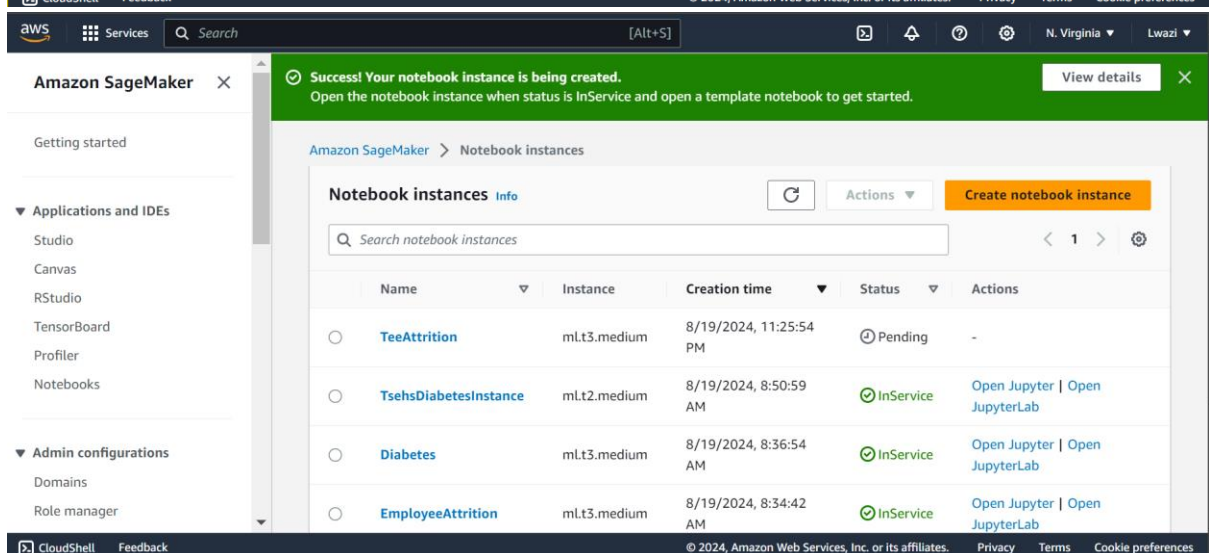
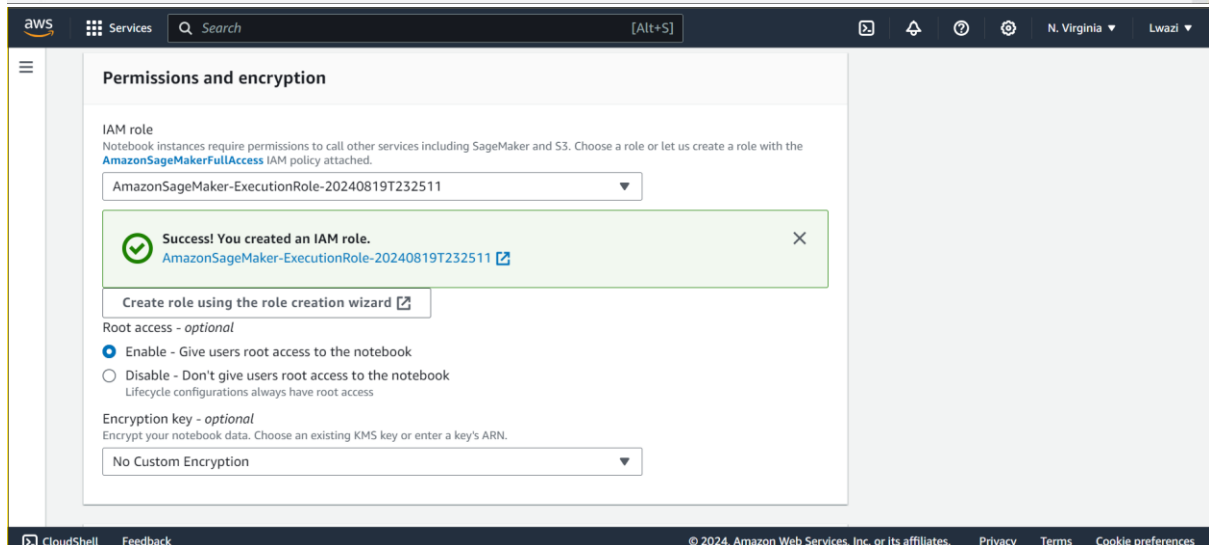
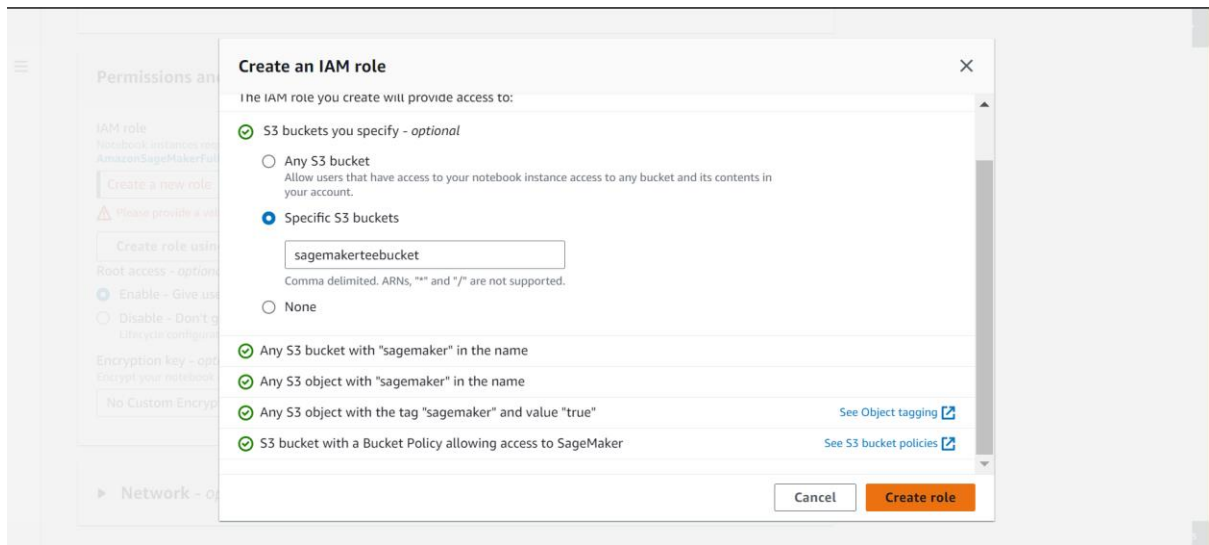
Feedback

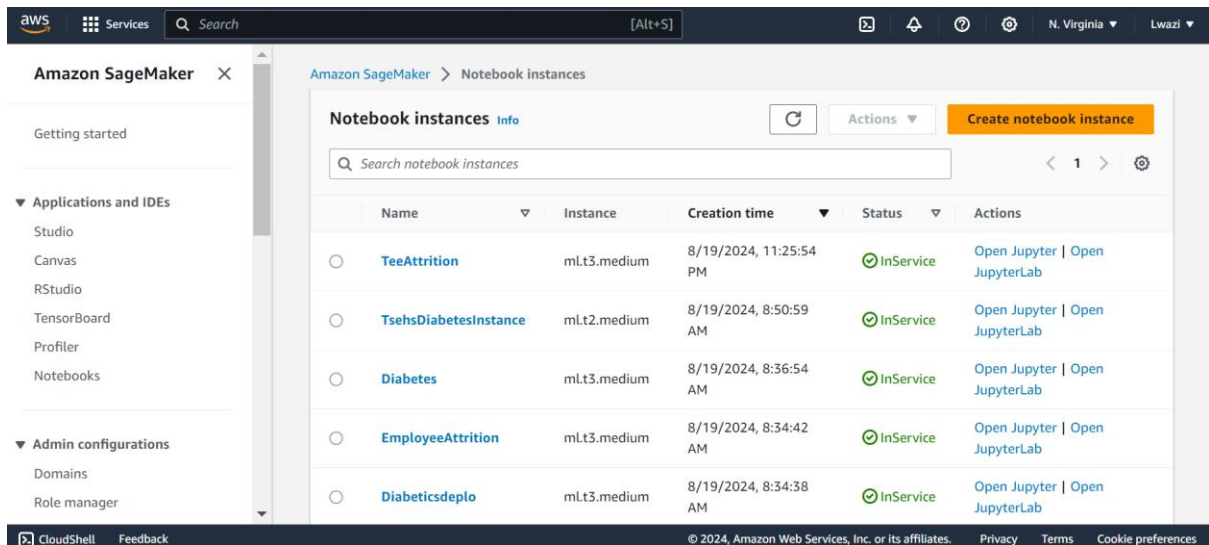
© 2024, Amazon Web Services, Inc. or its affiliates.

Privacy

Terms

Cookie preferences





Uploading the attrition dataset together with the google colab notebook to be edited:



Continuing from the data splitting portion of the project. The x and y, test and train data is combined, respectively, for better display.

Data Splitting

```
In [23]: #Libraries
from sklearn.model_selection import train_test_split
```

```
In [24]: x = data.drop(columns = ['Employee ID', 'Attrition'], axis=1)
y = data['Attrition']

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)
```

```
In [29]: #Displaying the training data
TrainData = x_train.join(y_train)
TrainData.head()
```

Out[29]:

	Age	Gender	Years at Company	Job Role	Monthly Income	Work-Life Balance	Job Satisfaction	Performance Rating	Number of Promotions	Overtime	Distance from Home	Education Level
2926	50	1	38	4	7026	2	1	0	1	0	3	3
3781	23	1	6	0	5002	0	0	0	2	1	14	1
4498	47	1	26	2	8037	2	3	0	0	0	35	1
4141	50	1	23	2	7252	2	2	3	0	0	96	2
5245	49	1	28	4	11314	1	0	0	0	0	14	1


```
In [30]: #Display all the test data
TestData=x_test.join(y_test)
TestData.head()
```

Out[30]:

	Age	Gender	Years at Company	Job Role	Monthly Income	Work-Life Balance	Job Satisfaction	Performance Rating	Number of Promotions	Overtime	Distance from Home	Education Level
5065	28	0	12	0	4638	1	2	0	0	0	61	1
5871	18	0	2	1	7769	1	3	0	0	0	39	1
3011	21	1	3	2	8030	3	2	1	2	0	71	4
3065	21	0	1	0	4268	2	2	0	0	1	54	0
1320	55	1	16	2	7527	2	1	0	1	0	10	3

```
In [37]: # Saving the trained data
TrainData.to_csv('TrainData.csv', index=False, index_label='Row', header=False)
```

```
In [38]: #Saving the test data
TestData.to_csv('TestData.csv', index=False, index_label='Row', header=False)
```



Open JupyterLab Quit Logout

Files Running Clusters Conda SageMaker Examples

Select items to perform actions on them.

Upload New ↕

<input type="checkbox"/>	0		Name	Last Modified	File size
<input type="checkbox"/>			Attrition.ipynb	Running a minute ago	910 kB
<input type="checkbox"/>			Attrition.csv	23 minutes ago	1.05 MB
<input type="checkbox"/>			TestData.csv	seconds ago	86.2 kB
<input type="checkbox"/>			TrainData.csv	a minute ago	345 kB

The saved data can be see on the notebook home list.

From there, the boto3 library is used to upload the files to the S3 bucket. The paths to the training and testing datasets are defined, their S3 paths are constructed, then they are uploaded to the to the designated S3 bucket loactions.

```

In [38]: #Saving the test data
TestData.to_csv('TestData.csv', index=False, index_label='Row', header=False)

In [ ]: #Uploading the trained and test data to Amazon S3 bucket

In [39]: #Necessary packages
import boto3
import re

In [40]: bucketName = 'sagemakerteebucket'
TrainFile = r'EmployeeAttritionData/TrainData/TrainData.csv'
TestFile = r'EmployeeAttritionData/TestData/TestData.csv'
ValFile = r'EmployeeAttritionData/ValData/Val.csv'
ModelFolder = r'EmployeeAttrition/model/'

In [41]: s3ModelOutput = r's3://{0}/{1}'.format(bucketName,ModelFolder)
s3Train = r's3://{0}/{1}'.format(bucketName,TrainFile)
s3Test = r's3://{0}/{1}'.format(bucketName,TestFile)
s3Val = r's3://{0}/{1}'.format(bucketName,ValFile)

In [42]: s3ModelOutput
Out[42]: 's3://sagemakerteebucket/EmployeeAttrition/model/'

In [43]: with open('TrainData.csv', 'rb') as f:
boto3.Session().resource('s3').Bucket(bucketName).Object(TrainFile).upload_fileobj(f)

In [44]: with open('TestData.csv', 'rb') as f:
boto3.Session().resource('s3').Bucket(bucketName).Object(TestFile).upload_fileobj(f)

```

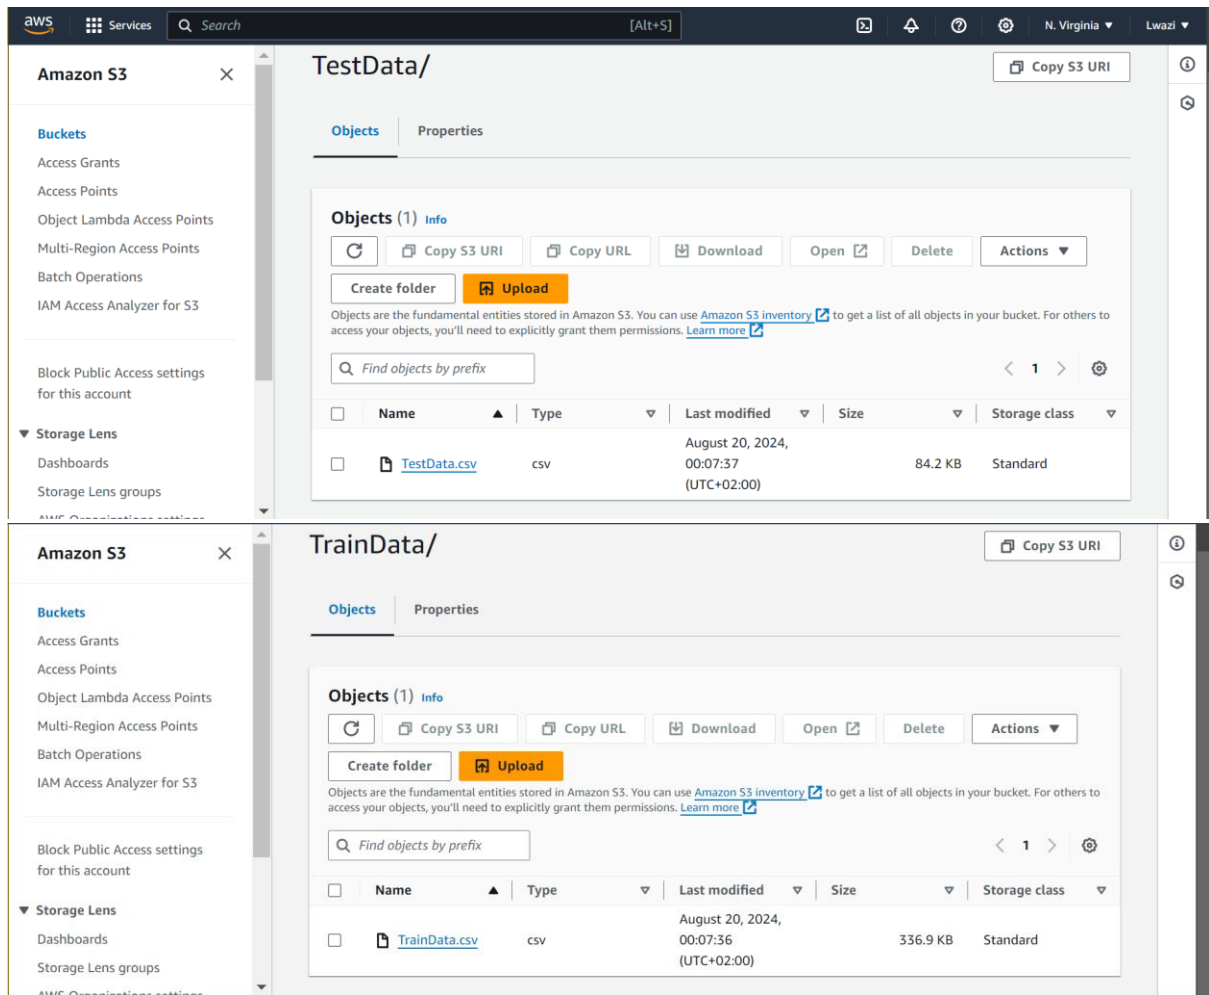
From the following snippets, it can be observed that the folders and datasets, have indeed been uploaded on the specified S3 bucket.

The first screenshot shows the AWS S3 console for the bucket 'sagemakerteebucket'. The 'Objects' tab is selected, showing a single object: 'EmployeeAttritionData/'. The object is a folder, and its details are shown below the table.

Name	Type	Last modified	Size	Storage class
EmployeeAttritionData/	Folder	-	-	-

The second screenshot shows the AWS S3 console for the folder 'EmployeeAttritionData/'. The 'Objects' tab is selected, showing two objects: 'TestData/' and 'TrainData/'. Both objects are folders, and their details are shown below the table.

Name	Type	Last modified	Size	Storage class
TestData/	Folder	-	-	-
TrainData/	Folder	-	-	-



To Be Continued...